

Capítulo 5. Deep learning para visión por computadora



5.1 Introducción a los convnets

Objetivo

- Introducir redes convolucionales o convnets
 - Clasificación de imágenes
 - *Data augmentation* para mitigar sobreajustes
 - Uso de convets pre-entrenados
 - Ajuste fino
 - Visualización del proceso de aprendizaje

MNIST: dígitos 0-9

- Utilizando DNN alcanzamos un *accuracy* del 97.8%
- Usaremos una pila de capas:
 - Conv2D
 - MaxPooling2D
- Entrada:
 - Tensor imagen: (altura, anchura, canales)
 - MNIST (28, 28, 1)
 - Un solo canal de grises

Sumario del modelo

- La salida de cada capa es un tensor
- (altura, anchura, canales)
- Dimensiones altura/anchura tienden a disminuir
- El número de canales está controlado por el 1er. argumento de las capas

```
from keras import layers
from keras import models

model = models.Sequential()
model.add(layers.Conv2D(32, (3, 3), \
                      activation='relu', \
                      input_shape=(28, 28, 1)))
model.add(layers.MaxPooling2D((2, 2)))
model.add(layers.Conv2D(64, (3, 3), \
                      activation='relu'))
model.add(layers.MaxPooling2D((2, 2)))
model.add(layers.Conv2D(64, (3, 3), \
                      activation='relu'))
```

- Añadir capas densas
 - Aplanar el modelo 3D → 1D: $(3, 3, 64) \rightarrow (576, 1)$
 - Añadir capas densas
 - Softmax para clasificación en última capa

```
model.add(layers.Flatten())
model.add(layers.Dense(64, activation='relu'))
model.add(layers.Dense(10, activation='softmax'))
```

- Cargar y preparar datos MNIST

- Entrenamiento y test

- Etiquetas

- Compilar y ajustar

- Tarda 2 min.

```
from keras.datasets import mnist
from keras.utils import to_categorical

(train_images, train_labels), \
(test_images, test_labels) = \
mnist.load_data()

train_images = train_images.reshape((60000, 28, 28, 1))
train_images = train_images.astype('float32') / 255

test_images = test_images.reshape((10000, 28, 28, 1))
test_images = test_images.astype('float32') / 255

train_labels = to_categorical(train_labels)
test_labels = to_categorical(test_labels)

model.compile(optimizer='rmsprop',
              loss='categorical_crossentropy',
              metrics=['accuracy'])
model.fit(train_images, train_labels, epochs=5, batch_size=64)
```

Evaluación

Modelo	Accuracy	Err. rate (1-Acc)	Factor reducción
DNN	0.978	0.022	--
Conv	0.993	0.007	0.682 ¹

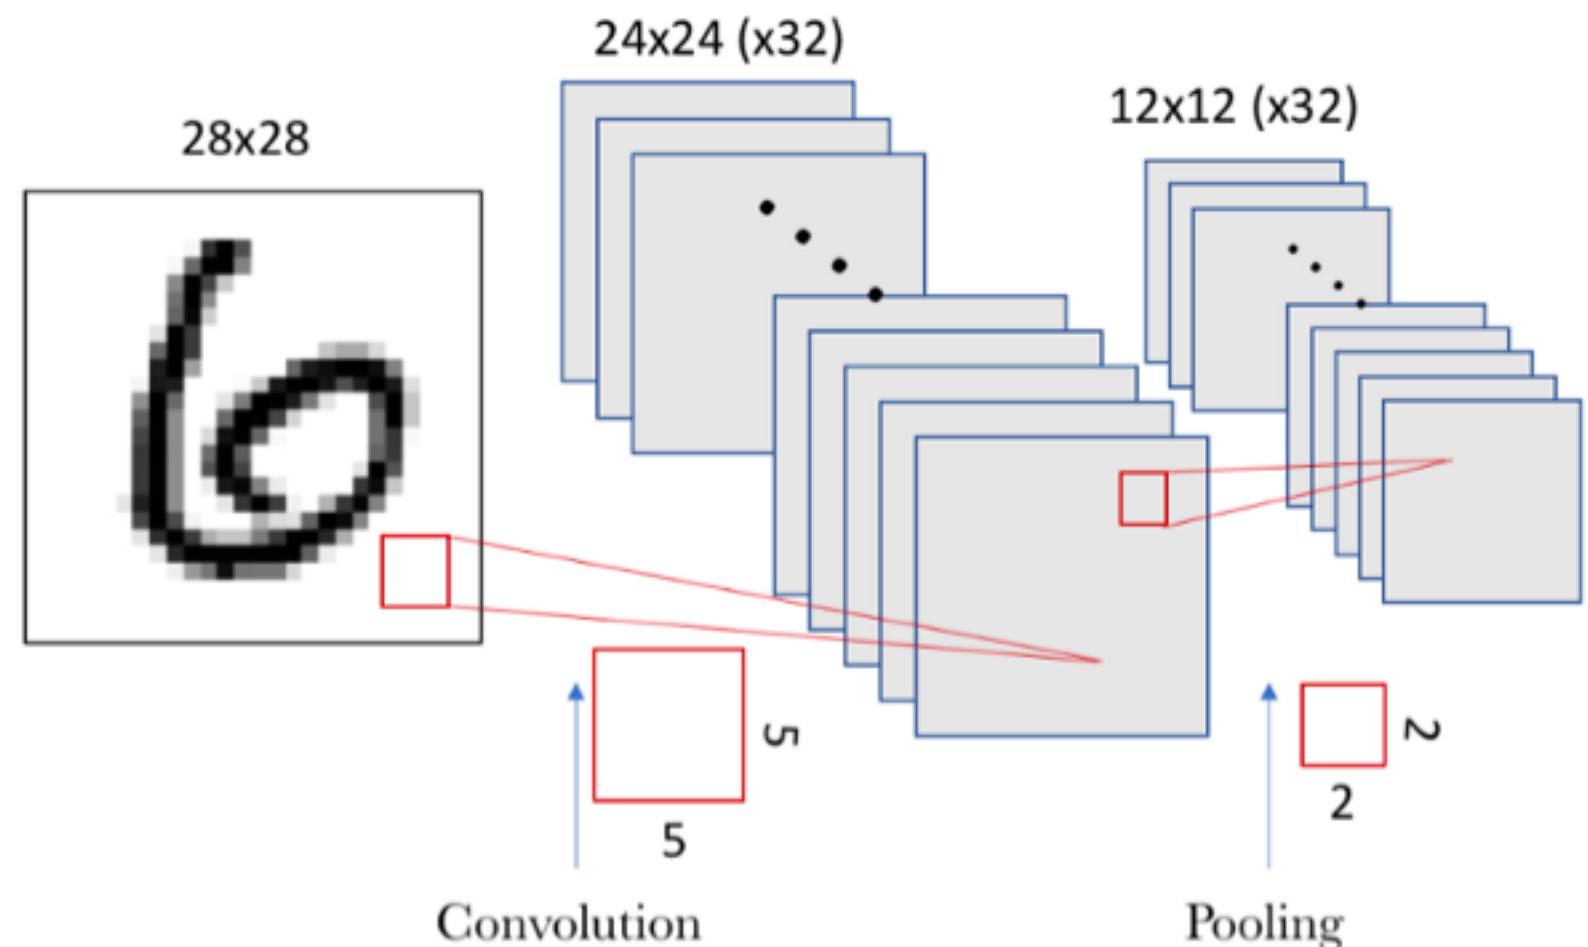
¹ Factor de reducción:

$$1 - \frac{Acc_{Conv}}{Acc_{DNN}} = 1 - \frac{0.007}{0.022} = 0.682$$

La operación de Convolución

Diferencias entre capas

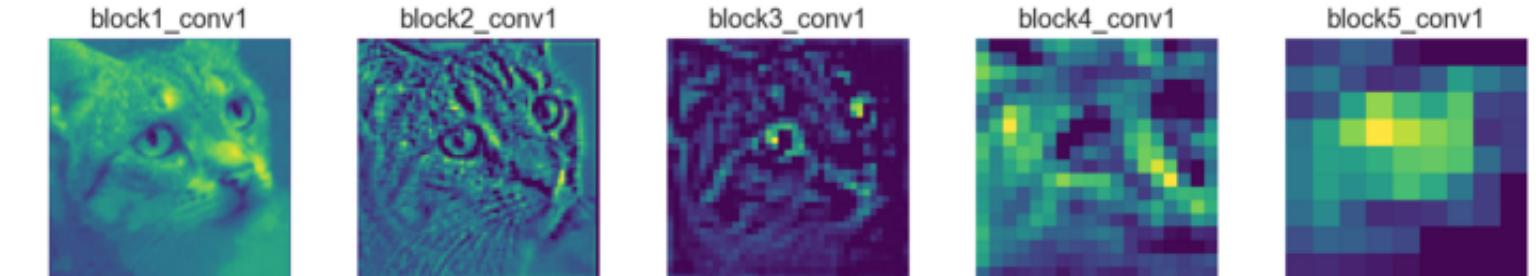
- Capas densas
 - Aprenden patrones globales
 - MINIST: todos los pixeles
 - Capas convolucionales²
 - Aprenden patrones locales
 - En imágenes: ventanas pequeñas



² Jordi Torres 2019, *Convolutional Neural Networks for beginners - Practical Guide with Python and Keras*

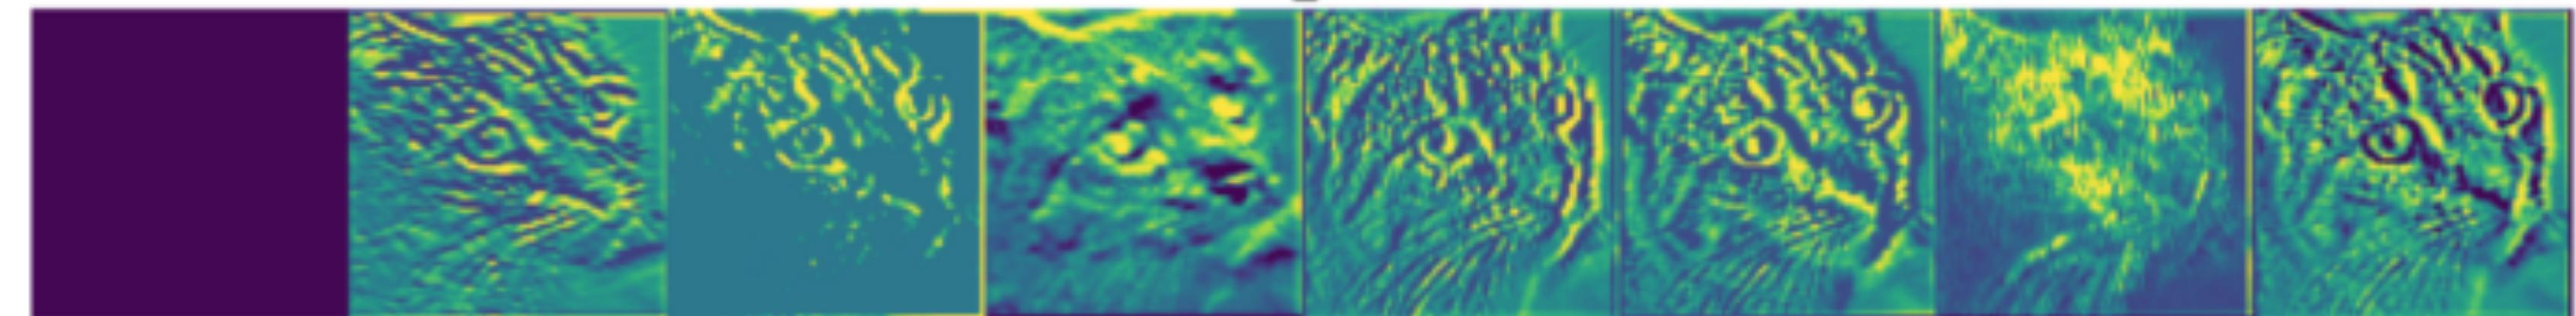
Propiedades de los convnets

- Los patrones son invariantes frente a traslaciones
- Aprenden jerarquías espaciales de patrones
- Patrones mayores y más abstractos en cada capa³

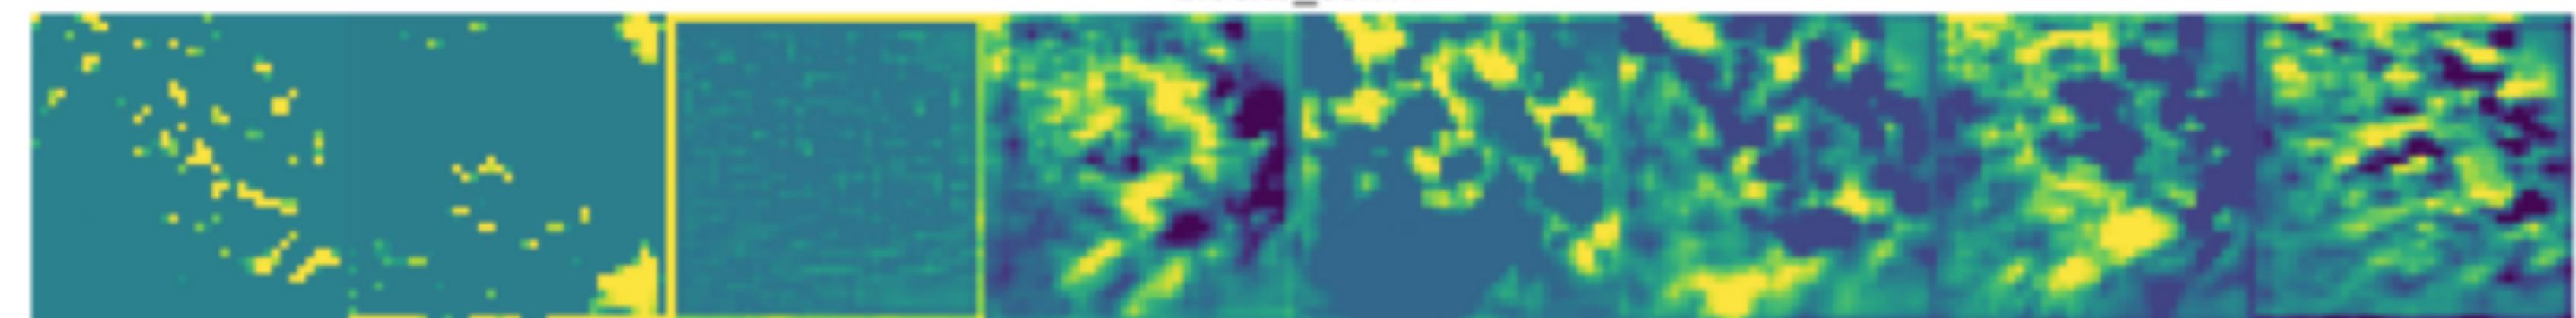


³ Arden Dertat, 2017, *Applied Deep Learning - Part 4: Convolutional Neural Networks*

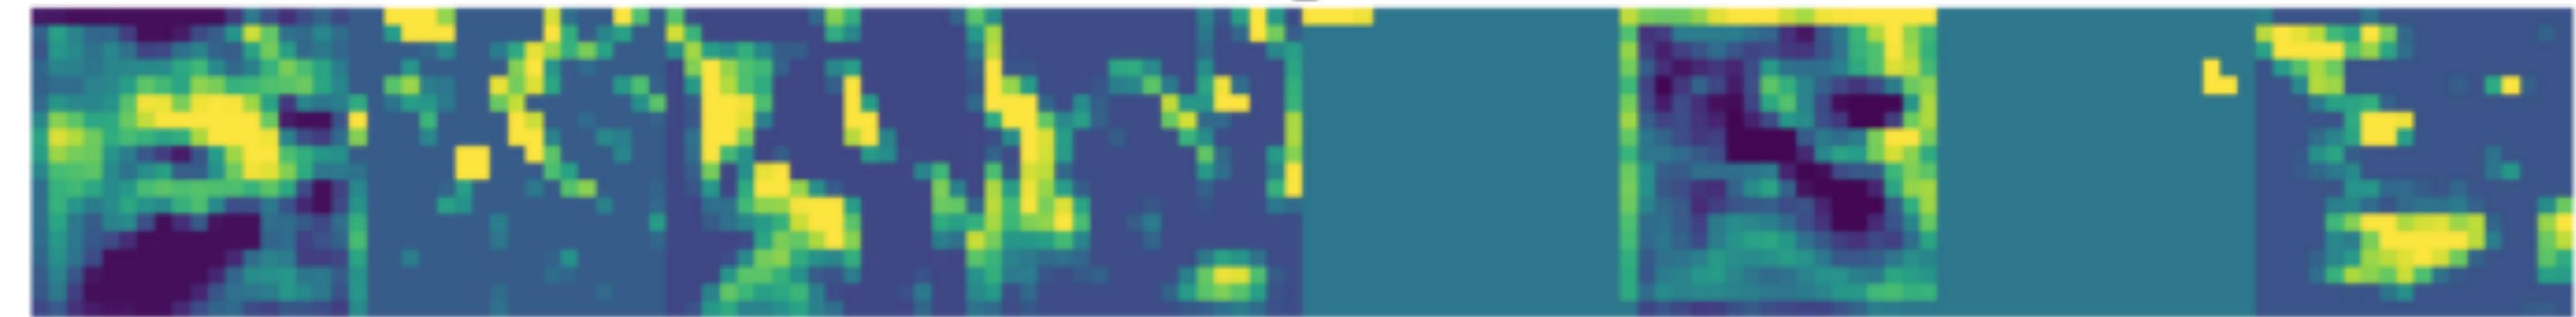
block2_conv1



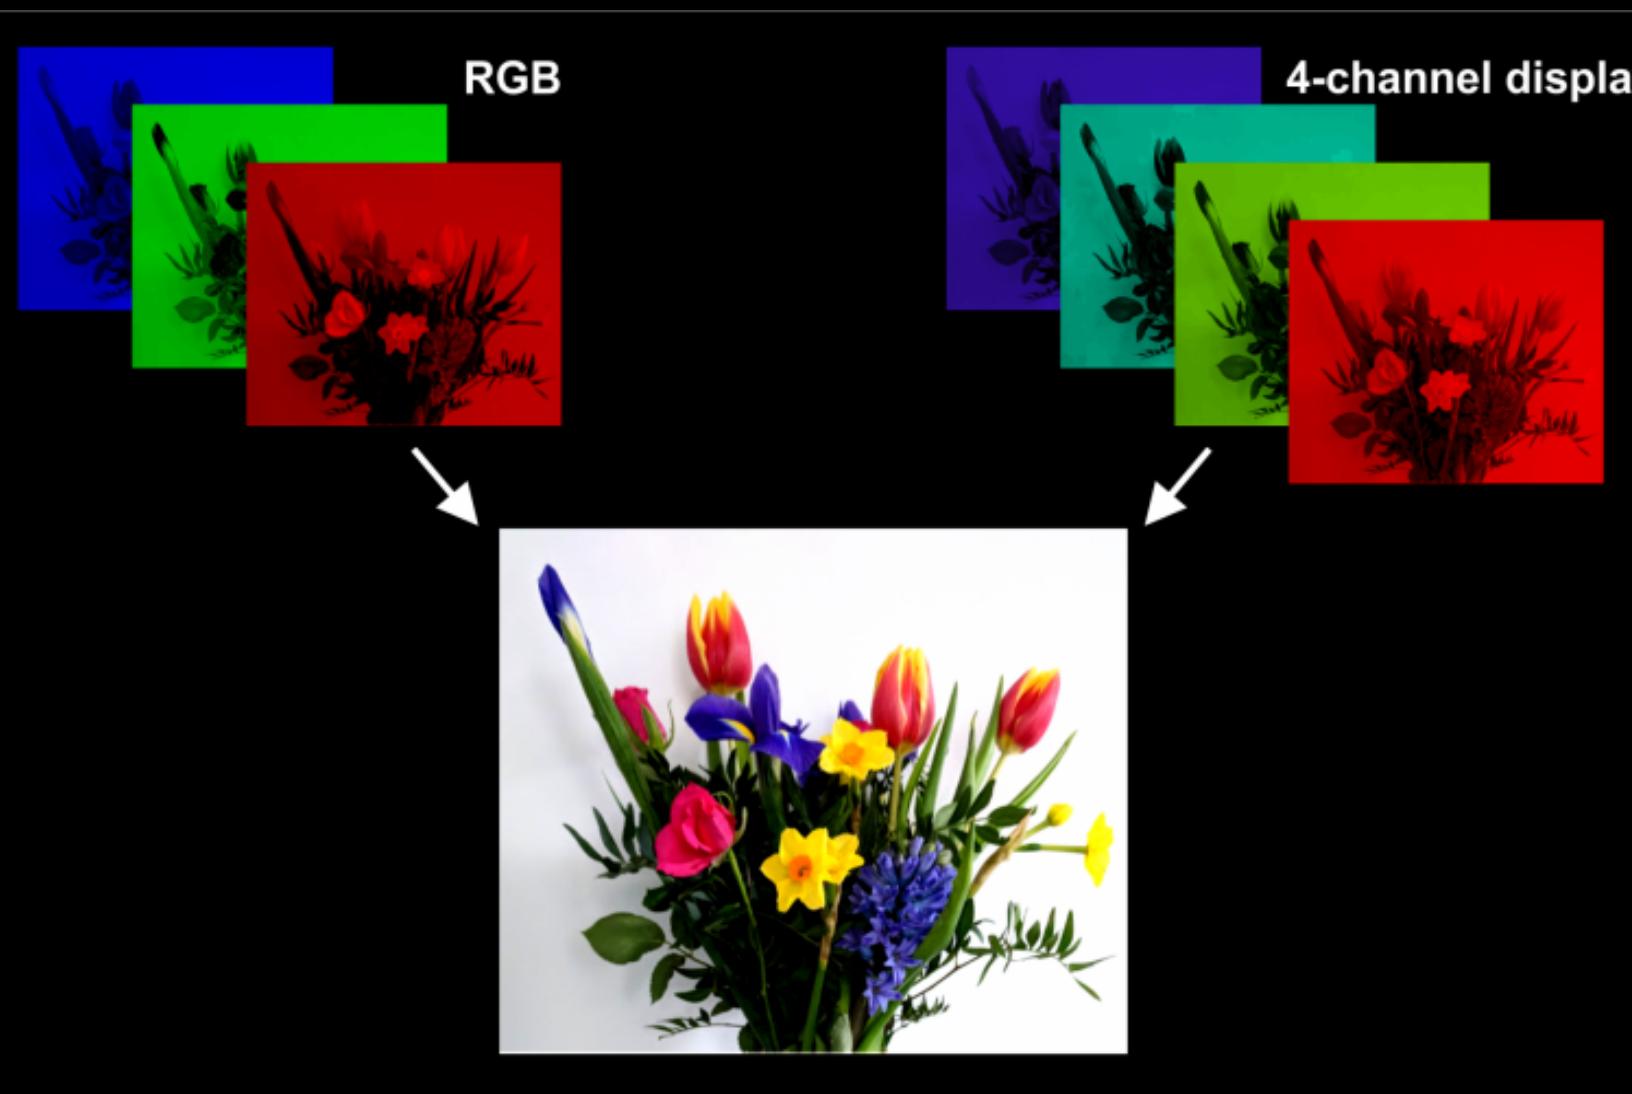
block3_conv1



block4_conv1



Feature maps (FM)



- Son los tensores 3D:
 - (Altura, Anchura, Profundidad / Canales)
 - Imagen RGB → 3 canales
 - Imagen B&W → 1 canal (niveles de grises)

Convolución:

- Extrae cuadros de los FM de entrada
- Transformaciones y FM de salida con tamaño FM arbitrarios
 - Primera capa: Profundidad son Canales de colores
 - Capas sucesivas: Profundidad son *filtros* que codifican aspectos de los datos⁴

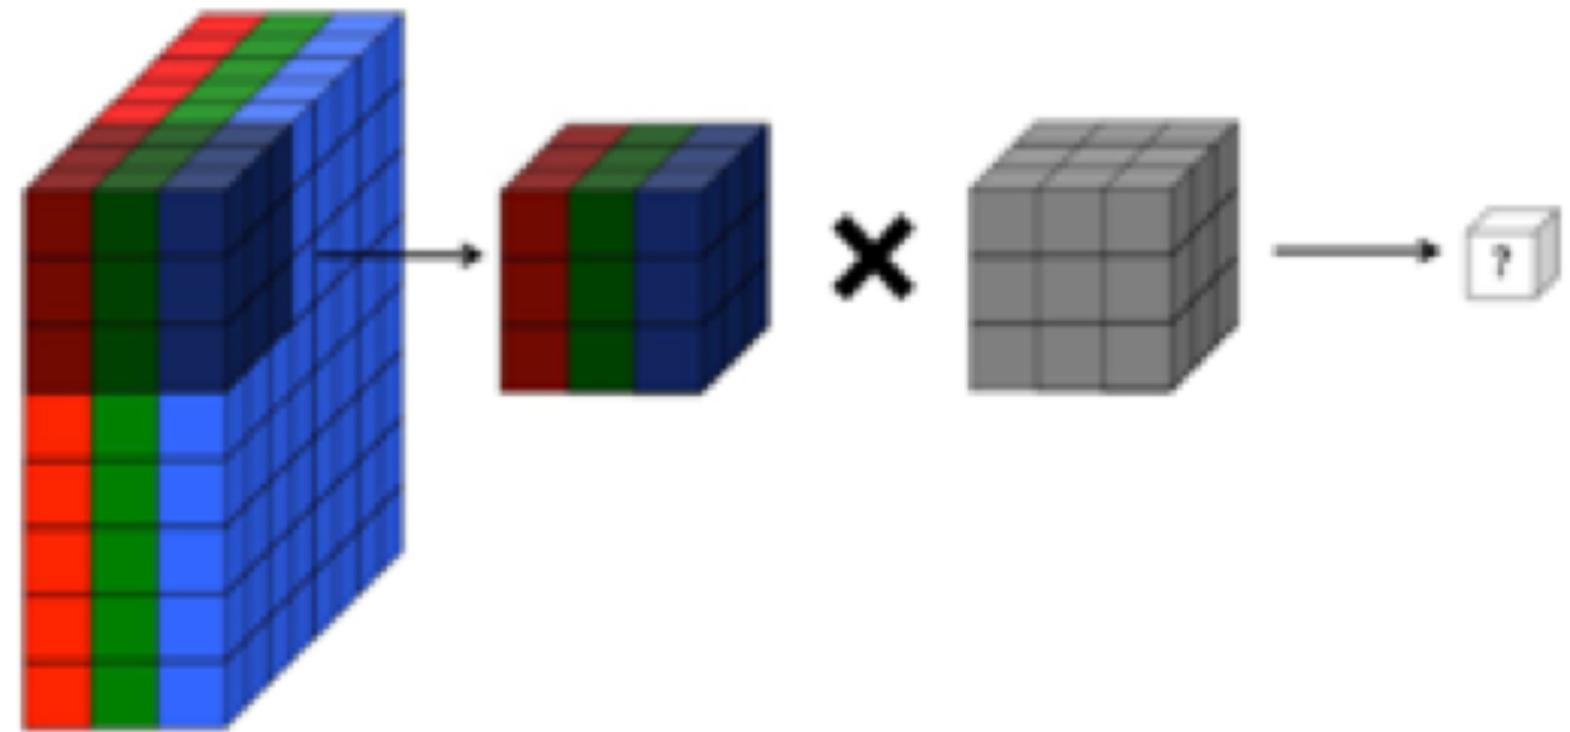


Figure 5-8. Representing a full-color RGB image as a volume and applying a volumetric convolutional filter

⁴ StackExchange, *In a CNN, does each new filter have different weights for each input channel, or are the same weights of each filter used across input channels?*

Efectos de borde

- Ejemplo⁵
 - FM 5x5
 - Ventana (filtro) 3x3
 - Sólo se pueden formar 9 cuadros completos

1	1	1	0	0
0	1	1	1	0
0	0	1	1	1
0	0	1	1	0
0	1	1	0	0

Input

1	0	1
0	1	0
1	0	1

Filter / Kernel

⁵ Arden Dertat, 2017, *Applied Deep Learning - Part 4: Convolutional Neural Networks*

1x1	1x0	1x1	0	0
0x0	1x1	1x0	1	0
0x1	0x0	1x1	1	1
0	0	1	1	0
0	1	1	0	0

Input x Filter

4		

Feature Map

1	1x1	1x0	0x1	0
0	1x0	1x1	1x0	0
0	0x1	1x0	1x1	1
0	0	1	1	0
0	1	1	0	0

Input x Filter

4	3	

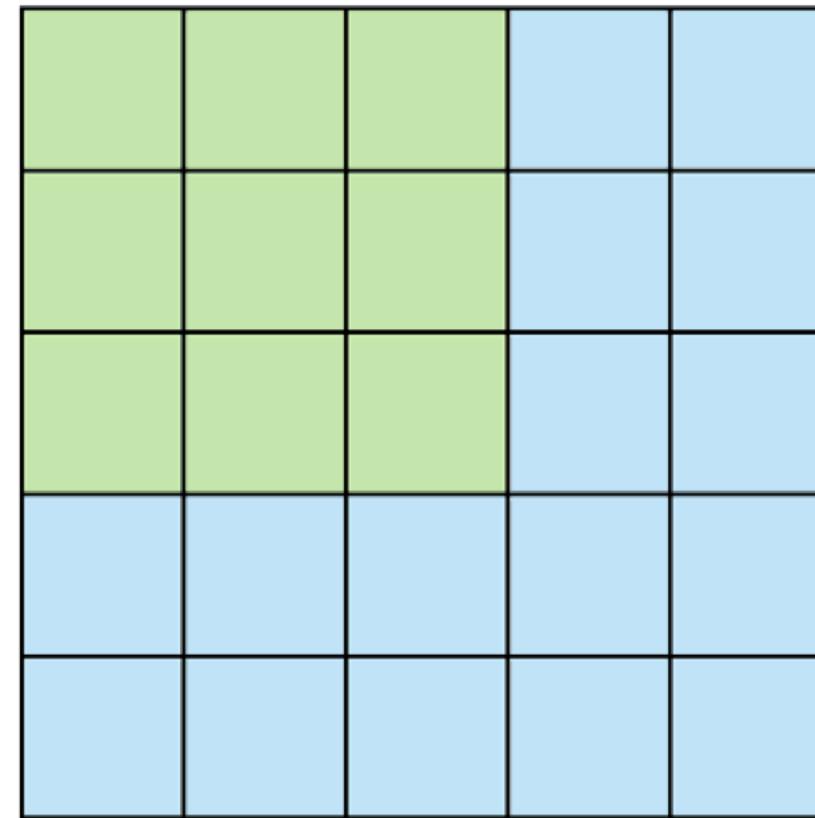
Feature Map

1x1	1x0	1x1	0	0
0x0	1x1	1x0	1	0
0x1	0x0	1x1	1	1
0	0	1	1	0
0	1	1	0	0

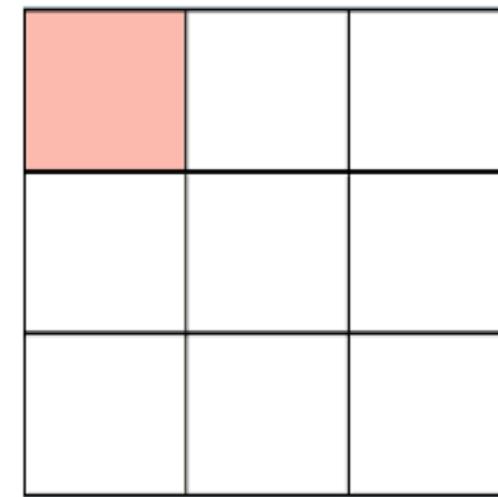
4		

Stride

- Pasos de convolución



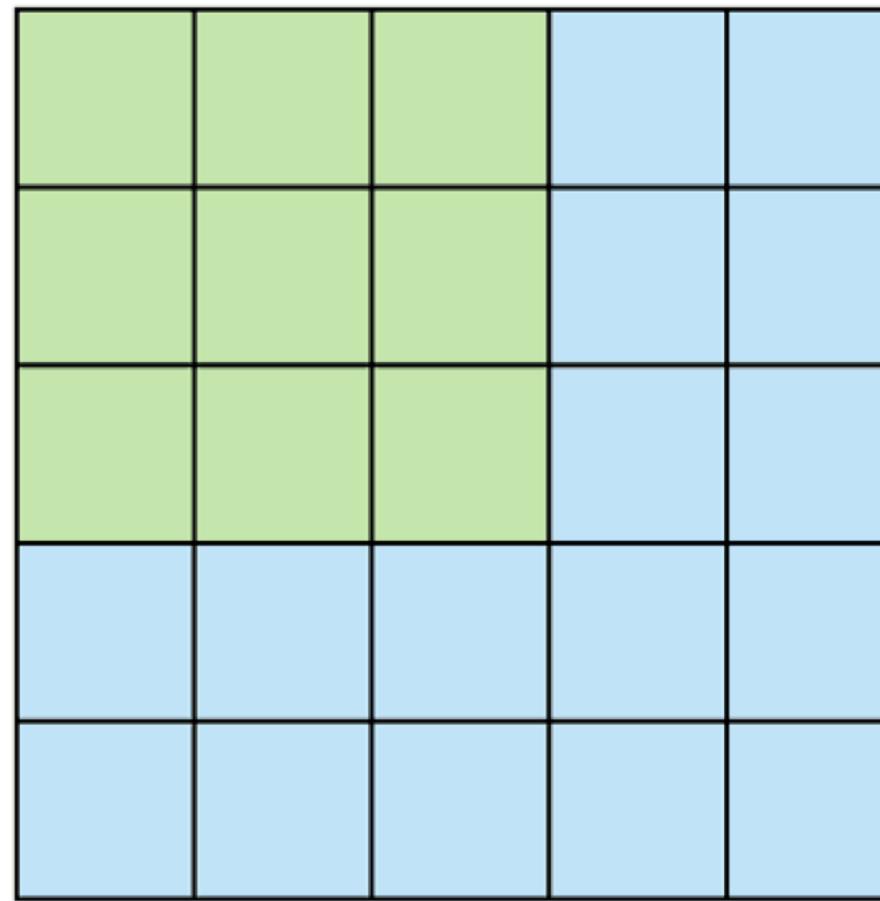
Stride 1



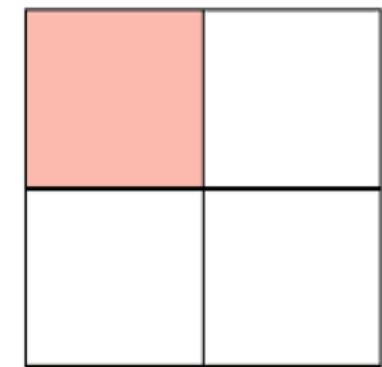
Feature Map

Stride

- Pasos de convolución
- Stride = 2 => Altura y anchura disminuyen por un factor 2
 - Además de efectos de borde



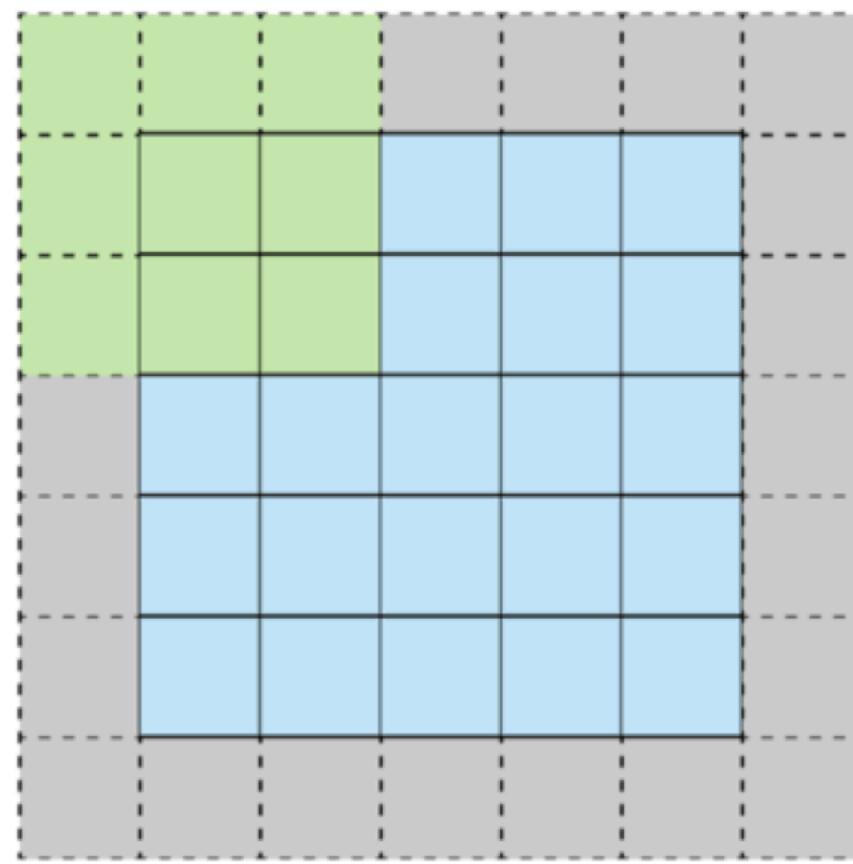
Stride 2



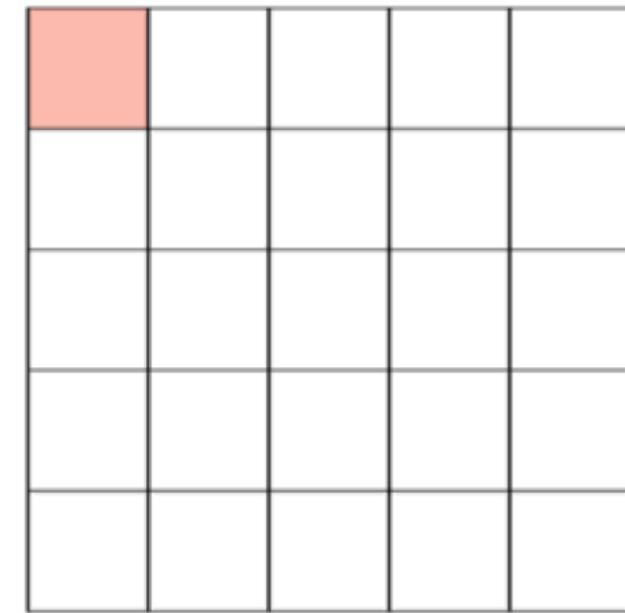
Feature Map

Padding

- Añadir artificialmente filas y columnas para mantener el tamaño de los FM entre capas
- Argumentos:
 - valid - **No** (sólo ventanas válidas)
 - same - **Sí**



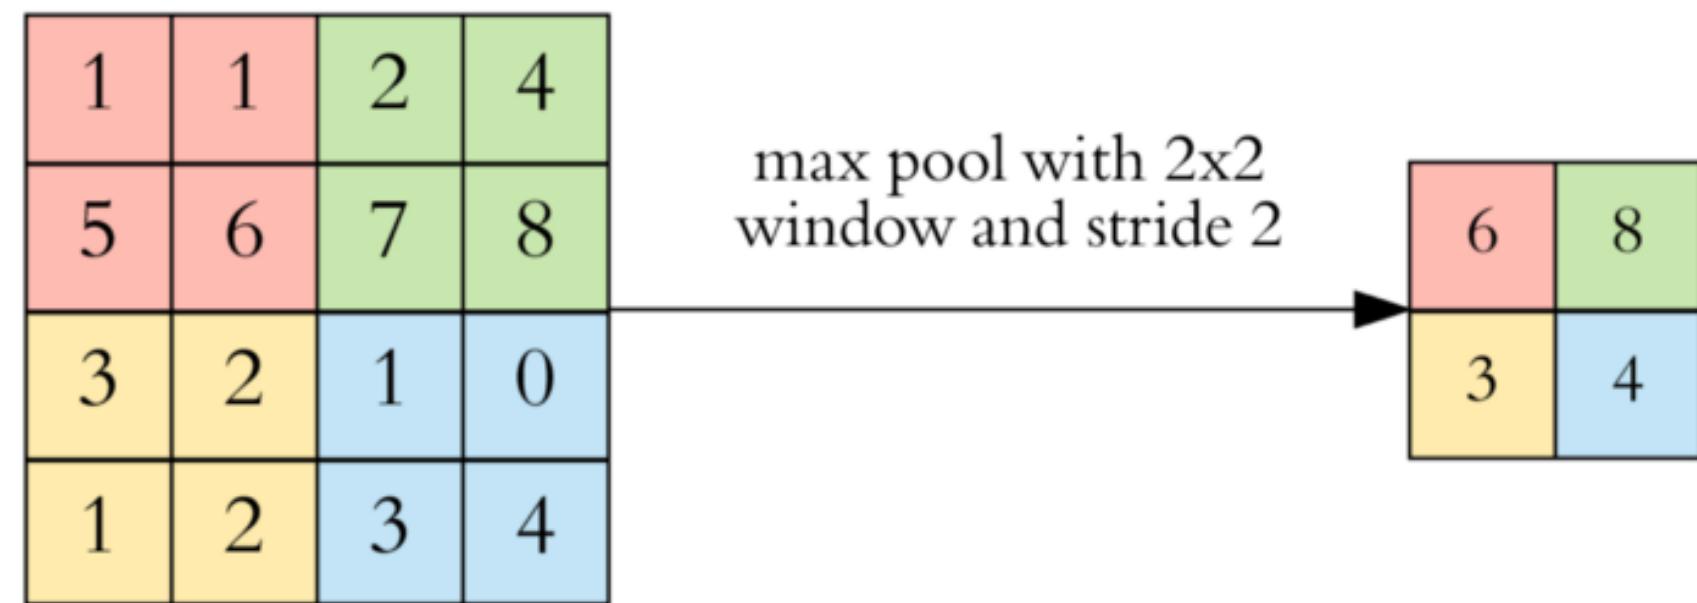
Stride 1 with Padding



Feature Map

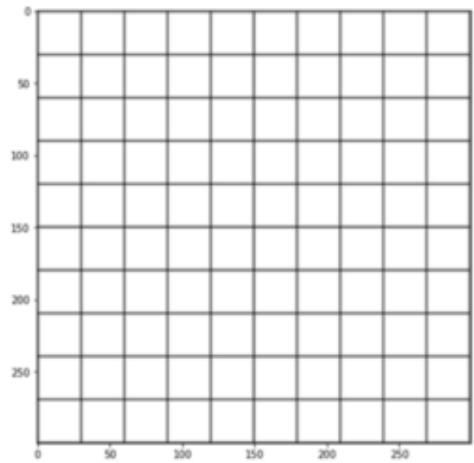
Max-pooling

- Necesarios para aprender la jerarquía y controlar el tamaño
- Disminuye el tamaño de los FM
- Extrae ventanas del FM con el valor máximo de cada canal
 - Aparte del *Max-pooling*, puede haber otros que extraigan promedio, mediana, etc.



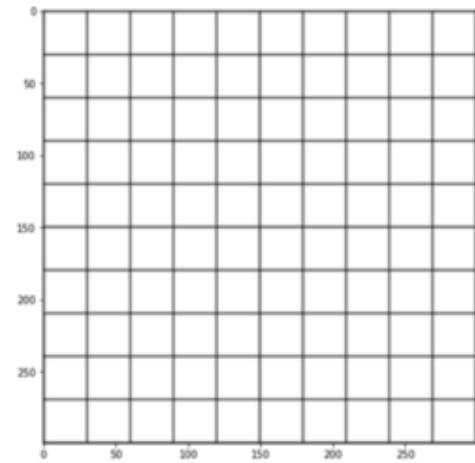
En arquitecturas de redes convolutivas, *pooling* (MaxPooling2D) se lleva a cabo con ventanas 2x2, *stride* 2 y sin *padding*, mientras que las convoluciones (Conv2D) se llevan a cabo en ventanas 3x3, *stride* 1 y con *padding*.

Ejemplos de filtros⁶



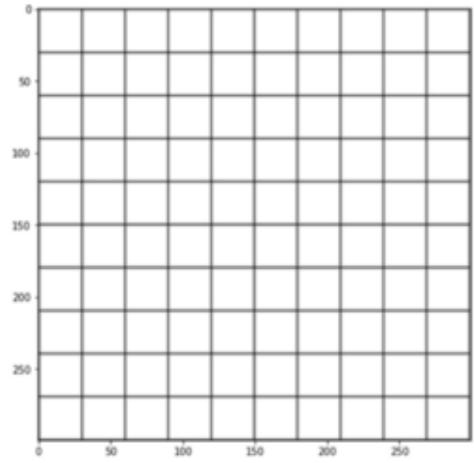
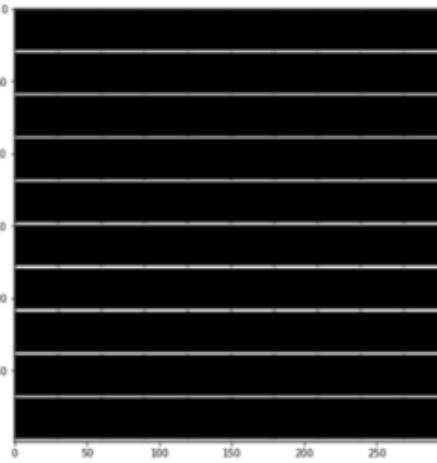
Naïve filter

0	0	0
0	1	0
0	0	0



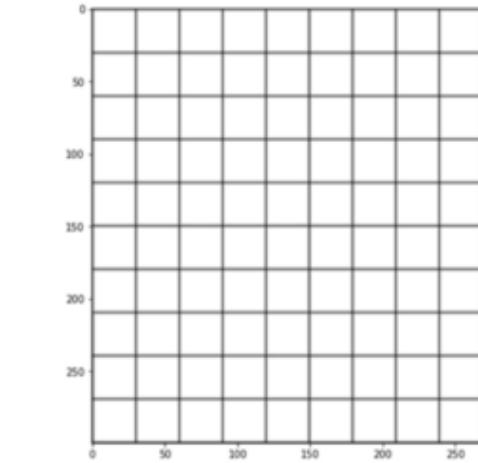
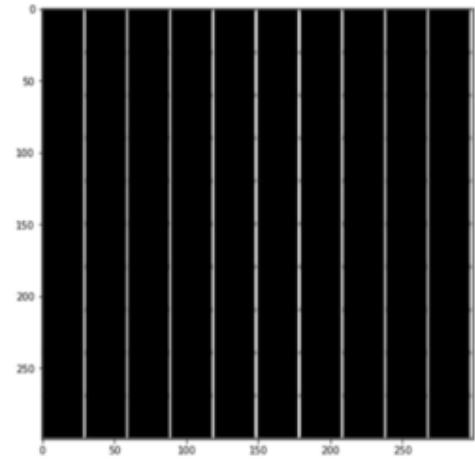
Horizontal edge filter

1	1	1
0	0	0
-1	-1	-1



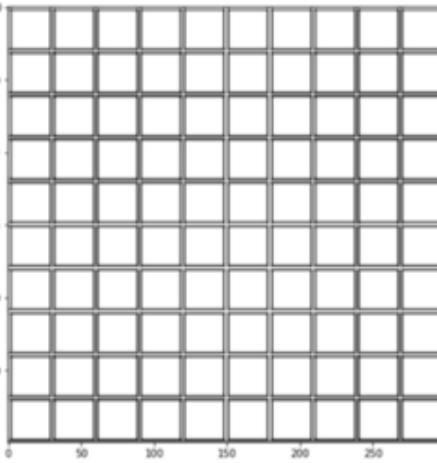
Vertical edge filter

1	0	-1
1	0	-1
1	0	-1



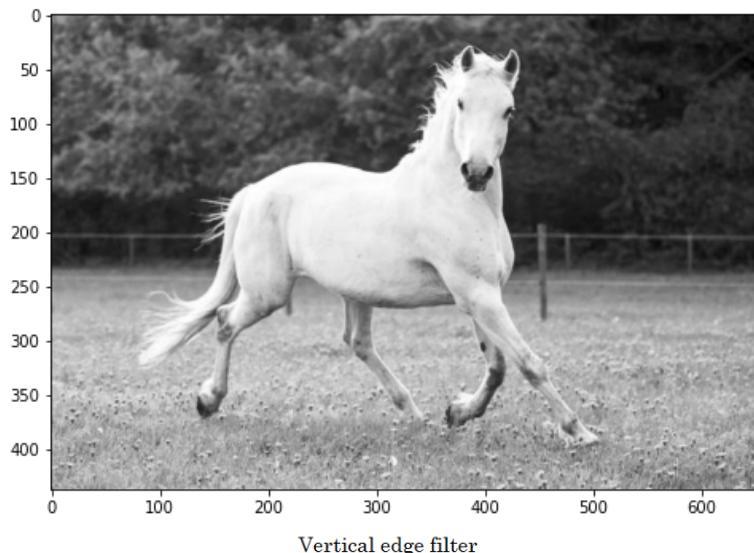
Edge filter

1	1	1
1	-7	1
1	1	1

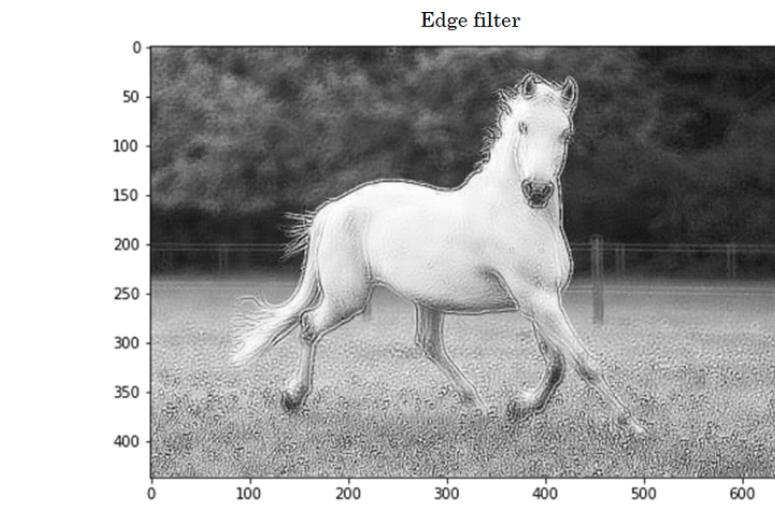
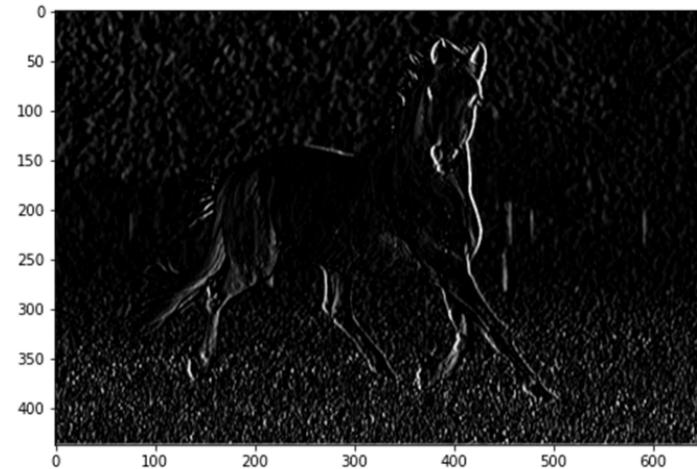


⁶ Fuente: Miguel Fernández Zafra, 2020, *Towards Data Science*

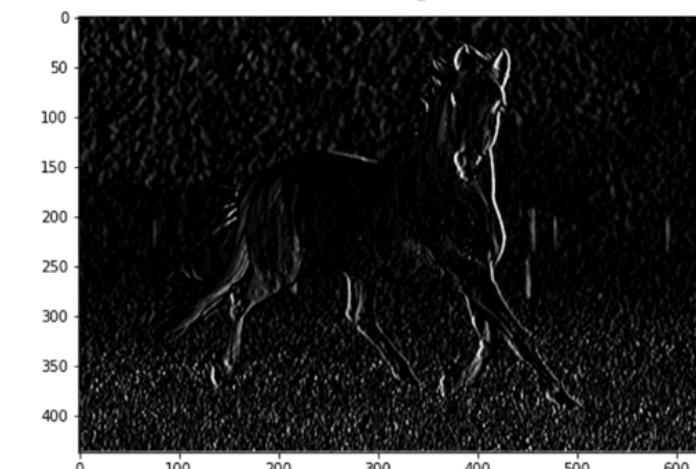
Ejemplos con fotografías reales⁶



Vertical edge filter



Edge filter



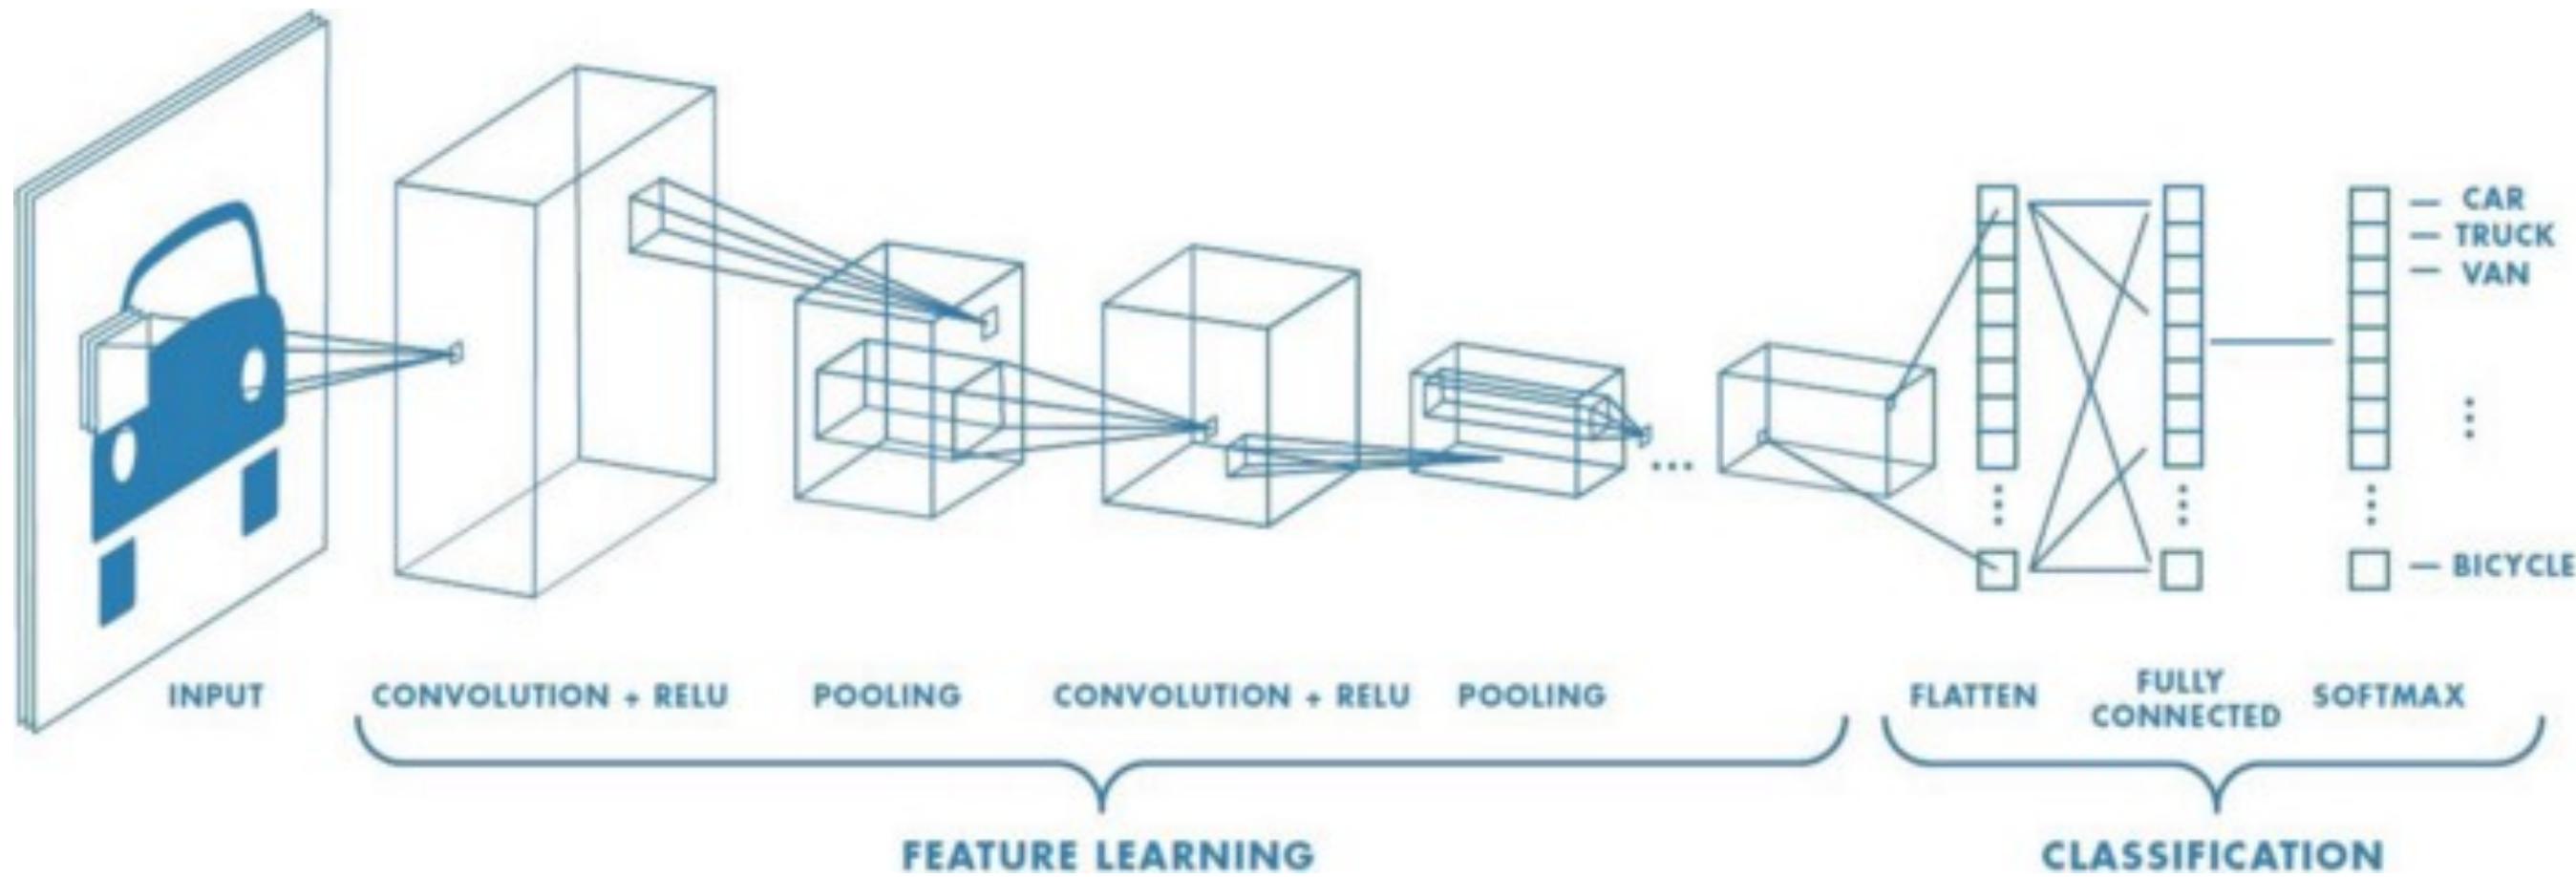
Vertical edge filter



2x2 Max Pooling

⁶ Fuente: [Miguel Fernández Zafra, 2020, Towards Data Science](#)

Representación de red convolutiva



© Sumit Saha 2018, Towards Data Science: [A Comprehensive Guide to Convolutional Neural Networks – the ELI5 way](#)

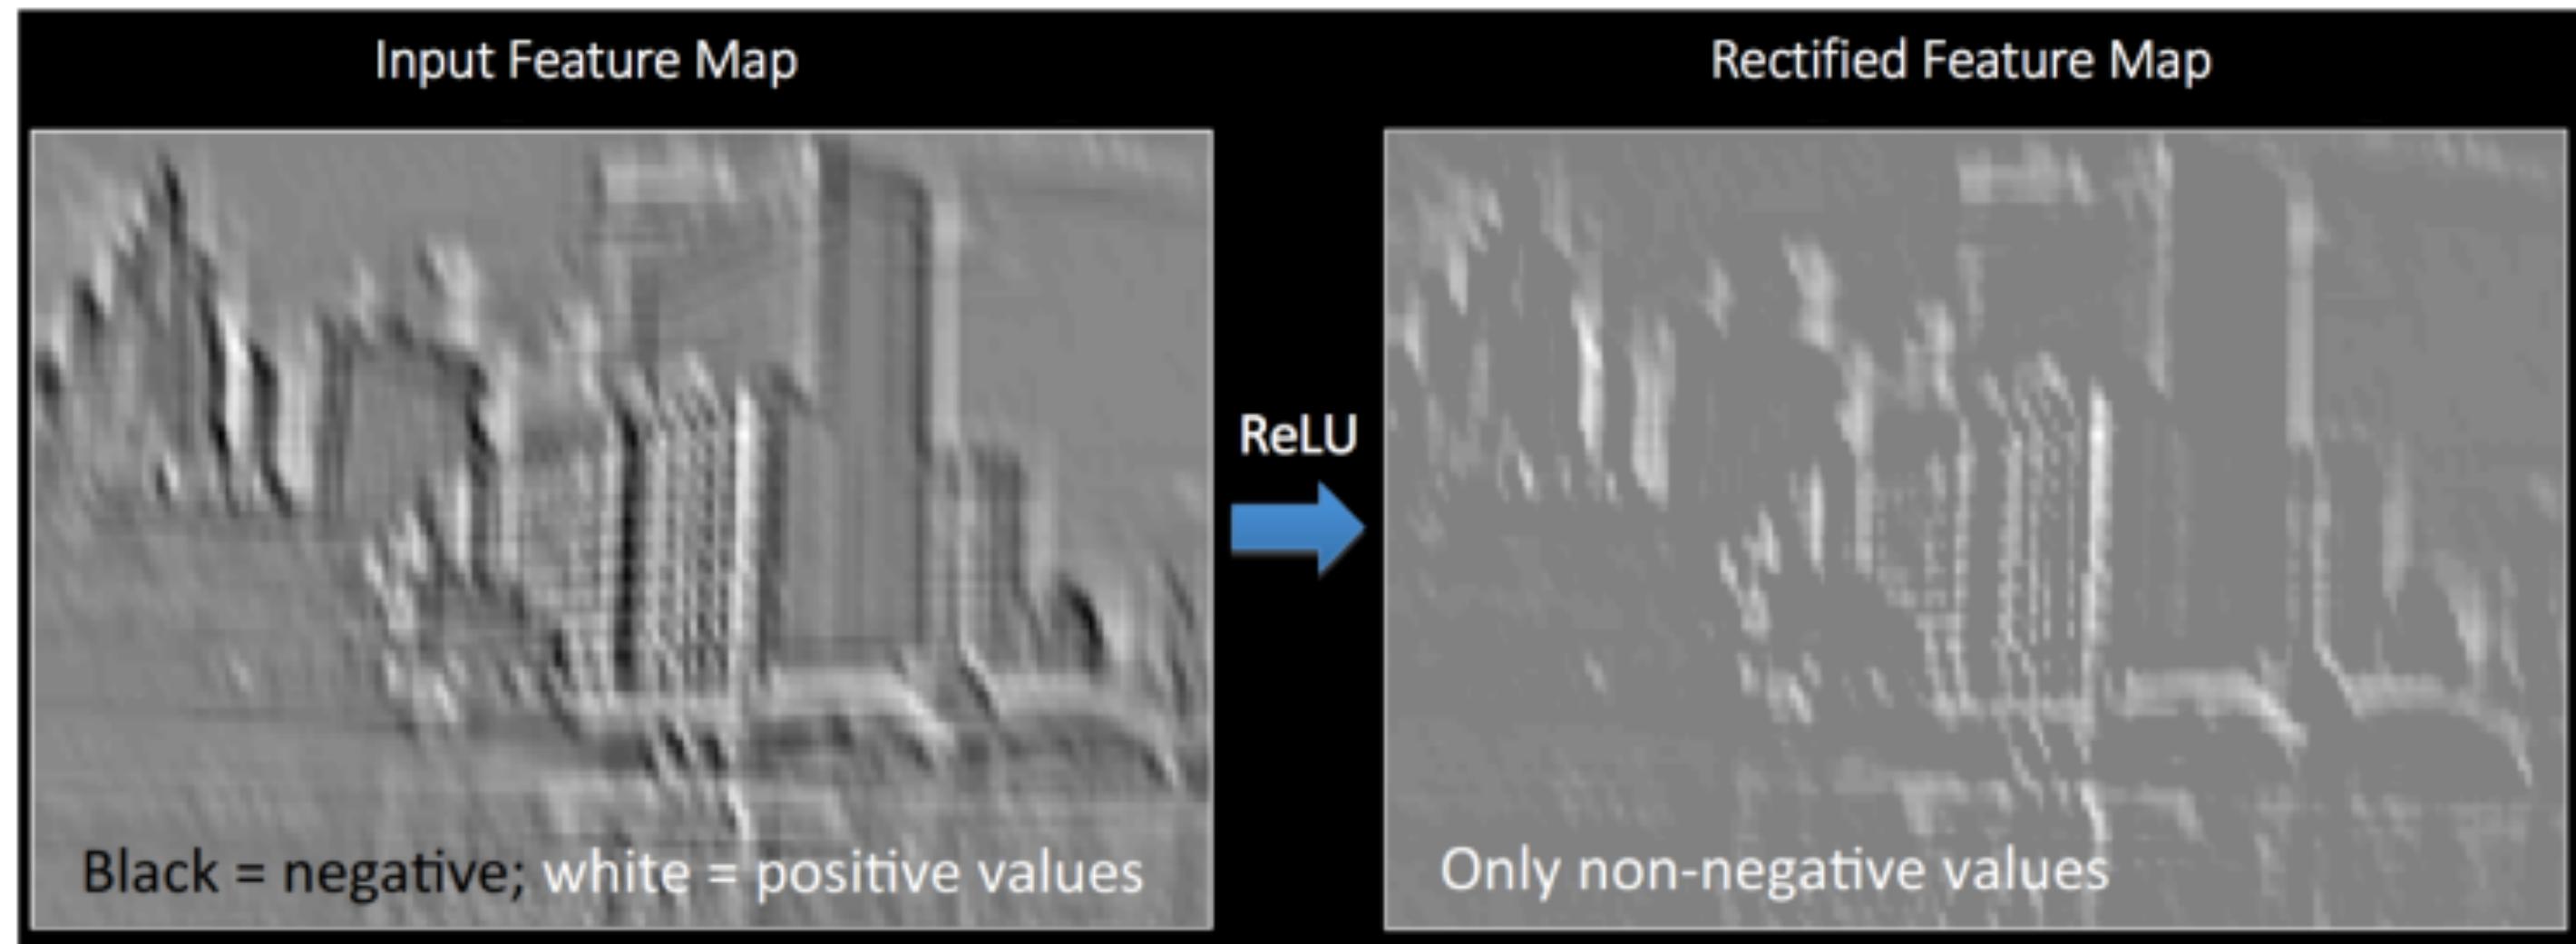
Filtros

jcgonzalez, 2017, aps: Use of convolutional neural network
for image classification

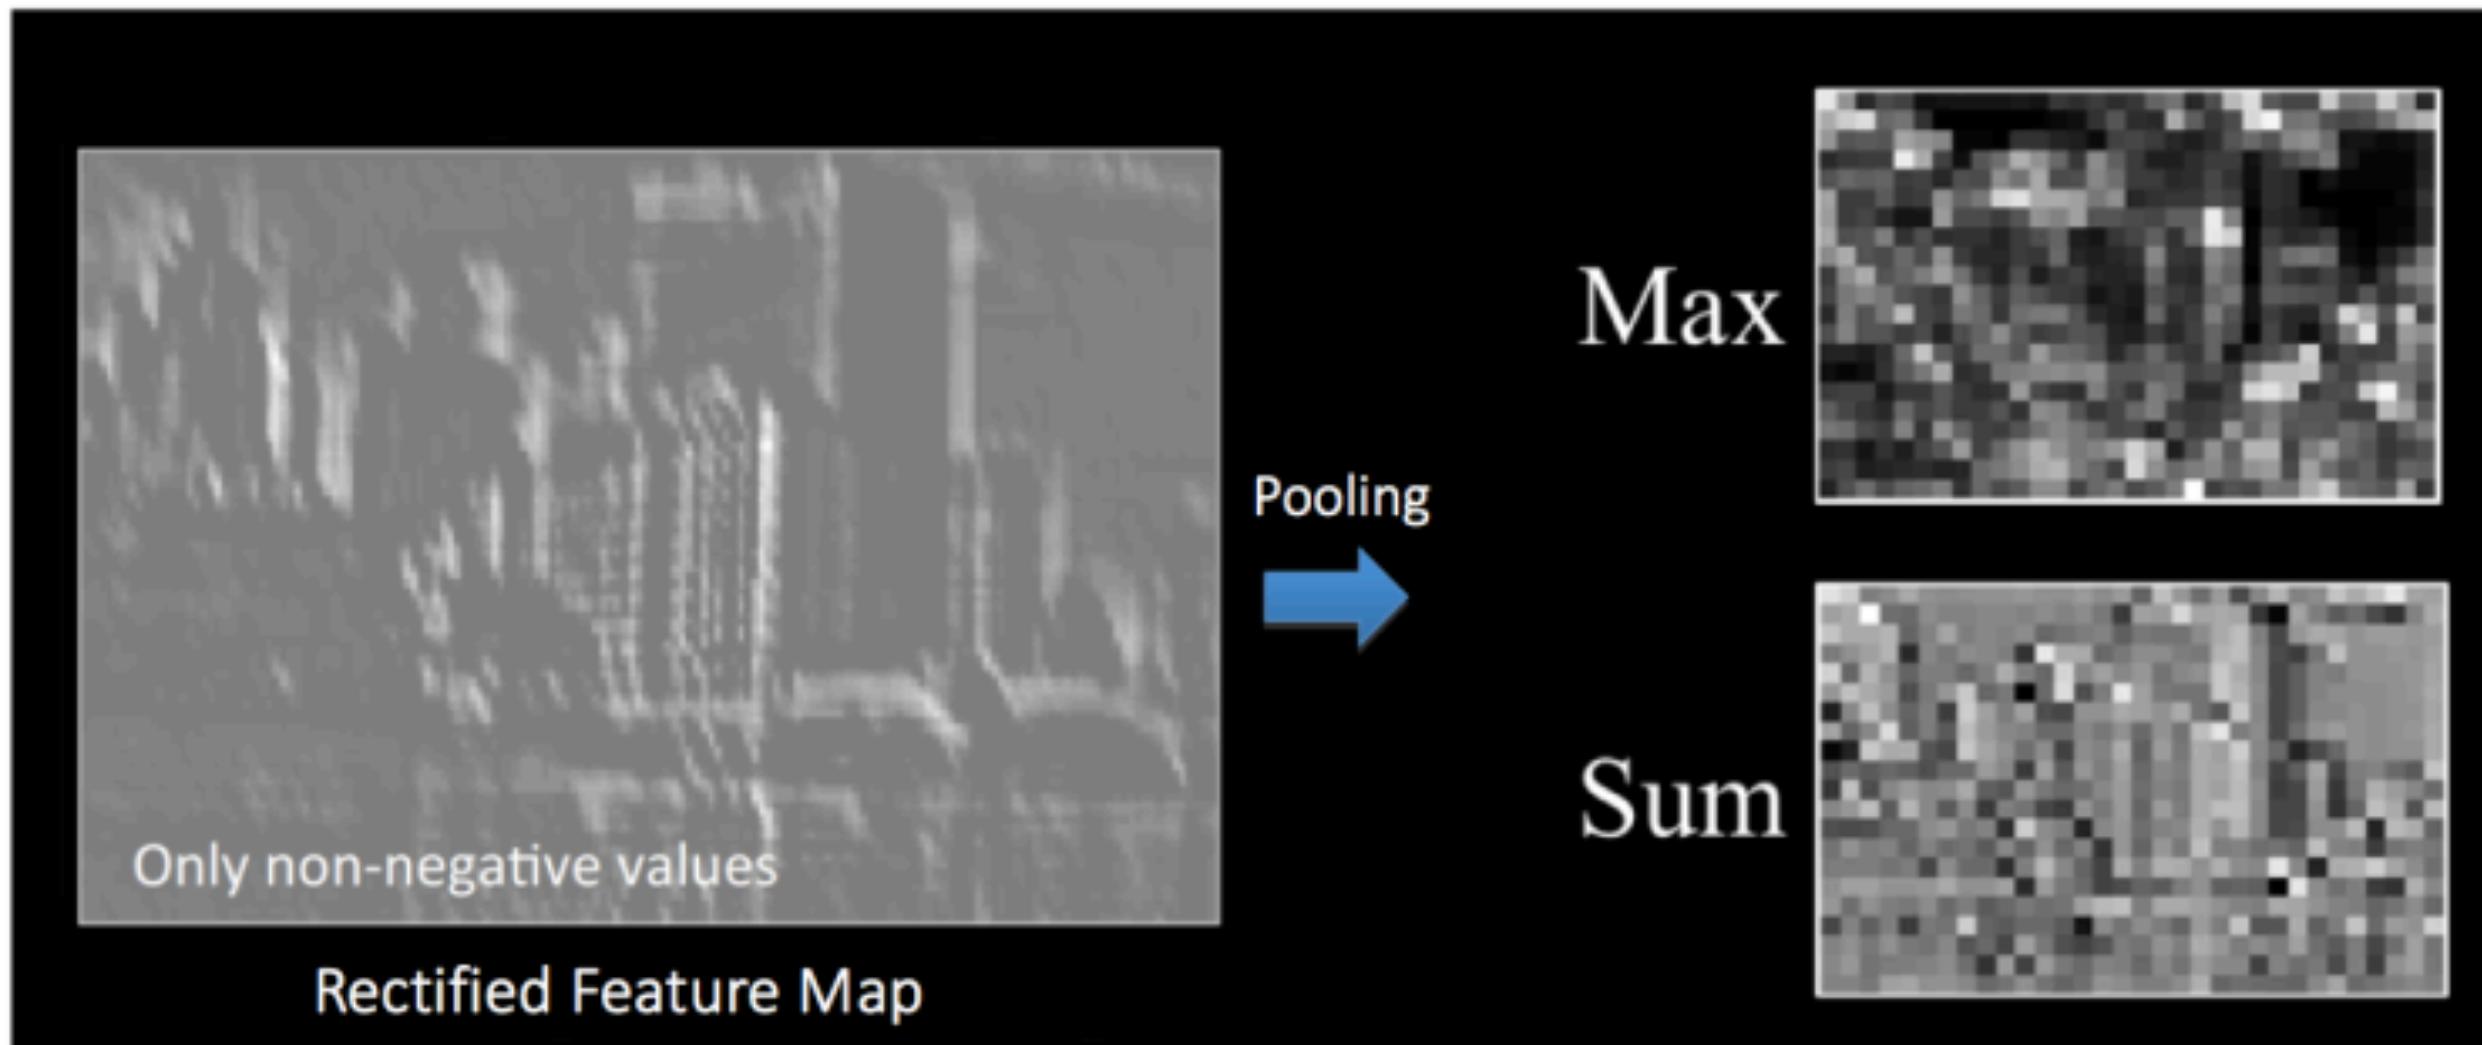


Input

Activación ReLu



Max y Sum Pooling



*Convnet con base
de datos pequeña*

El próximo día



KEEP
CALM
AND
SEE YOU
NEXT CLASS